

# Markov decision process

## Definition

A Markov decision process (MDP)  $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$  consists of

- ▶  $\mathcal{S}$  a set of states (in my case  $\subseteq \mathbb{R}^n$ )
- ▶  $\mathcal{A}$  a set of actions (in my case *finite*)
- ▶  $P : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{P}(\mathcal{S})$  its Markov transition kernel
- ▶  $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{P}(\mathbb{R})$  its immediate reward distribution
- ▶  $\gamma \in (0, 1)$  the discount factor

# Q-Learning

- ▶ Policy:  $\pi : \mathcal{S} \rightarrow \mathcal{P}(\mathcal{A})$
- ▶ State-value function:  $V^\pi : \mathcal{S} \rightarrow \mathbb{R}$   
$$V^\pi(s) = \mathbb{E} \left( \sum_{t \geq 0} \gamma^t R_t \mid R_t \sim R(S_t, A_t), S_t \sim P(S_{t-1}, A_{t-1}), A_t \sim \pi(S_t), S_0 = s \right)$$
- ▶ State-action-value (Q-) function  $Q^\pi : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$

$$Q^\pi(s, a) = \mathbb{E}(R(s, a) + \gamma V^\pi(S_0) \mid S_0 \sim P(s, a))$$

- ▶ Optimal Q-function is defined as

$$Q^*(s, a) = \sup_{\pi} Q^\pi(s, a)$$

- ▶ One can show that there is a policy  $\pi^*$  such that  $Q^* = Q^{\pi^*}$ .  
This is the optimal policy - the goal of Q-learning (and Reinforcement Learning in general).

# Artificial Neural Networks

## Definition

An **ANN** (Artificial Neural Network) with structure  $\{d_i\}_{i=0}^{L+1} \subseteq \mathbb{N}$ , activation functions  $\sigma_i = (\sigma_{ij} : \mathbb{R} \rightarrow \mathbb{R})_{j=1}^{d_i}$  and weights  $\{W_i \in M^{d_i \times d_{i-1}}, v_i \in \mathbb{R}^{d_i}\}_{i=1}^{L+1}$  is the function  $F : \mathbb{R}^{d_0} \rightarrow \mathbb{R}^{d_{L+1}}$

$$F(x) = w_{L+1} \circ \sigma_L \circ w_L \circ \sigma_{L-1} \circ \cdots \circ w_1 x$$

where  $w_i$  is the affine function  $x \mapsto W_i x + v_i$  for  $i \in [L+1]$ .

Here  $\sigma_i(x_1, \dots, x_{d_i}) = (\sigma_{i1}(x_1), \dots, \sigma_{id_i}(x_{d_i}))$ .

$L \in \mathbb{N}_0$  is called the number of hidden layers.

$d_i$  is the number of neurons or nodes in layer  $i$ .

An ANN is called *deep* if there are two or more hidden layers.

# The Bellman operator

Denote  $\pi_Q$  as the *greedy* policy with respect to  $Q$  i.e.  $\pi(s, a) = 1$  for  $a = \operatorname{argmax}_a Q(s, a)$ . For every policy  $\pi$  we define the operators

$$(P^\pi Q)(s, a) = \mathbb{E}(Q(S', A') \mid S' \sim P(s, a), A' \sim \pi(S'))$$

$$(T^\pi Q)(s, a) = \mathbb{E}R(s, a) + \gamma(P^\pi Q)(s, a)$$

$T^\pi$  is called the Bellman operator. It can be shown that  $Q^\pi$  is a fixed point for  $T^\pi$ . Finally we define *Bellmans optimality operator*  $T$  as

$$TQ = T^{\pi_Q} Q$$

Bellmans optimality equation is then  $TQ^* = Q^*$ .

# Context

## Theorem

*If both  $\mathcal{S}$  and  $\mathcal{A}$  are finite, and  $R$  is deterministic, then the simple iteration*

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha_t(s_t, a_t)[R_t + \gamma Q_t(s_{t+1}, \pi_{Q_t}(s_{t+1})) - Q_t(s_t, a_t)]$$

*converges with probability 1 to  $Q^*$ , given that*

$$\sum_{t \geq 1} \alpha_t(s, a) = \infty, \quad \sum_{t \geq 1} \alpha_t^2(s, a) < \infty$$

*for all  $(s, a) \in \mathcal{S} \times \mathcal{A}$ .*

## Theorem (Universal approximation theorem)

*An ANN with 1 hidden layer is sufficient to approximate any continuous function  $[0, 1]^k \rightarrow \mathbb{R}$  (at cost of layer-width).*

# Sparse ReLU Network

## Definition

For  $s, V \in \mathbb{R}$  a  $(s, V)$ -**Sparse ReLU Network** is an ANN  $f$  with any structure  $\{d_i\}_{i \in [L+1]}$ , all activation functions being *ReLU* i.e.  $\sigma_{ij} = \max(\cdot, 0)$  and any weights  $(W_\ell, v_\ell)$  satisfying

- ▶  $\max_{\ell \in [L+1]} \|\widetilde{W}_\ell\|_\infty \leq 1$
- ▶  $\sum_{\ell=1}^{L+1} \|\widetilde{W}_\ell\|_0 \leq s$
- ▶  $\max_{j \in [d_{L+1}]} \|f_j\|_\infty \leq V$

Here  $\widetilde{W}_\ell = (W_\ell, v_\ell)$ .

The set of them we denote  $\mathcal{F}(s, V)$ .

# The algorithm

---

**Algorithm 1:** Fitted Q-Iteration Algorithm

---

**Input:** MDP  $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$ , function class  $\mathcal{F}$ , sampling distribution  $\nu$ , number of iterations  $K$ , number of samples  $n$ , initial estimator  $\tilde{Q}_0$

**for**  $k = 0, 1, 2, \dots, K - 1$  **do**

    Sample i.i.d. observations  $\{(S_i, A_i), i \in [n]\}$  from  $\nu$  obtain

$R_i \sim R(S_i, A_i)$  and  $S'_i \sim P(S_i, A_i)$

    Let  $Y_i = R_i + \gamma \cdot \max_{a \in \mathcal{A}} \tilde{Q}_k(S'_i, a)$

    Update action-value function:

$$\tilde{Q}_{k+1} \leftarrow \operatorname{argmin}_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n (Y_i - f(S_i, A_i))^2$$

Define  $\pi_K$  as the greedy policy w.r.t.  $\tilde{Q}_K$

**Output:** An estimator  $\tilde{Q}_K$  of  $Q^*$  and policy  $\pi_K$

---

# The theorem

For any  $K \in \mathbb{N}$  let  $Q^{\pi_K}$  be the action-value function corresponding to policy  $\pi_K$  which is return by Algorithm 1, when run with a sparse ReLU network on the form

$$\mathcal{F}_0 = \{f(\cdot, a) \in \mathcal{F}(L^*, \{d_j^*\}_{j=0}^{L^*+1}, s^*) \mid a \in \mathcal{A}\}$$

where

$$L^* \lesssim (\log n)^{\xi'}, d_0 = r, d_j^*, d_{L+1} = 1, \lesssim n^{\xi'}, s^* \asymp n^{\alpha^*} \cdot (\log n)^{\xi'}$$

Let  $\mu$  be any distribution over  $\mathcal{S} \times \mathcal{A}$ . Under *some assumptions*

$$\|Q^* - Q^{\pi_K}\|_{1,\mu} \leq C \cdot \frac{\phi_{\mu,\nu} \cdot \gamma}{(1-\gamma)^2} \cdot |\mathcal{A}| \cdot (\log n)^{\xi^*} \cdot n^{(\alpha^*-1)/2} + \frac{4\gamma^{K+1}}{(1-\gamma)^2} \cdot R_{\max}$$

Here  $\alpha^* \in (0, 1)$ ,  $C, \xi', \xi^*, \phi_{\mu,\nu} \in \mathbb{R}_+$  are constants depending on the assumptions and  $R_{\max}$  the maximum possible reward.



# Naive interpretation

This paper shows that for *any* MDP satisfying mild(?) assumptions, its optimal policy  $\pi^*$  can be approximated arbitrarily close by the fitted Q-iteration algorithm (in  $\|\cdot\|_1$  over the chosen  $\mu$  distribution). I.e. this algorithm 'solves' (up to any precision) a large class of games, decision processes, etc.



# Problems

- ▶ Are the assumptions actually mild?
- ▶ How large are the constants / how quick is convergence?
- ▶ In the algorithm we assumed that we could solve a least squares problem on a function class of Neural Networks with several restrictions. According to one reviewer this is an NP-hard problem.
- ▶ Our result depended on a distribution  $\mu$ , so it does not say much about how close we are to  $\pi^*$  outside the support of  $\mu$ .
- ▶ The fitted Q-iteration algorithm differs from normal Deep Q-Learning in two important ways:
  - ▶ It avoids analysing errors in SGD and Back-Propagation by assuming that a global optimum is found.
  - ▶ It uses a fixed distribution on  $\mathcal{S} \times \mathcal{A}$  for batch sampling during *experience replay* rather than picking uniformly from actual experiences.

# Hölder Smoothness

## Definition

Let  $\mathcal{D} \subseteq \mathbb{R}^r$  be compact and  $\beta, H > 0$ . A function  $f : \mathcal{D} \rightarrow \mathbb{R}$  we call Hölder smooth if

$$\sum_{\alpha: |\alpha| < \beta} \|\partial^\alpha f\|_\infty + \sum_{\alpha: \|\alpha\|_1 = \lfloor \beta \rfloor} \sup_{x \neq y} \frac{|\partial^\alpha (f(x) - f(y))|}{\|x - y\|_\infty^{\beta - \lfloor \beta \rfloor}} \leq H$$

Where  $\alpha = (\alpha_1, \dots, \alpha_r) \in \mathbb{N}^r$ . We write  $f \in C_r(\mathcal{D}, \beta, H)$ .

We consider families of *Compositions of Hölder Functions*

$$\mathcal{G}(\{p_j, t_j, \beta_j, H_j\}_{j \in [q]})$$

where  $t_j, p_j \in \mathbb{N}$ ,  $t_j \leq p_j$  and  $H_j, \beta_j > 0$ , defined as containing  $f$  when  $f = g_q \circ \dots \circ g_1$  for  $g_j : [a_j, b_j]^{p_j} \rightarrow [a_{j+1}, b_{j+1}]^{p_{j+1}}$  functions on some real hypercubes that only depend on  $t_j$  of their inputs for each of their components  $g_{jk}$ , and satisfies  $g_{jk} \in C_{t_j}([a_j, b_j]_j^{t_j}, \beta_j, H_j)$ .

# Assumption 1, Hölder smoothness of $\mathcal{F}_0$ under $T$

Let

$$\mathcal{G}_0 = \{f : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R} : f(\cdot, a) \in \mathcal{G}(\{p_j, t_j, \beta_j, H_j\}_{j \in [q]}), \forall a \in \mathcal{A}\}$$

It is assumed that  $Tf \in \mathcal{G}_0$  for any  $f \in \mathcal{F}_0$ .

I.e. when using the Bellman optimality operator on our sparse ReLU networks, we should stay in the class of compositions of Hölder smooth functions.

## Assumption 2, Concentration Coefficients

Let  $\nu_1, \nu_2 \in \mathcal{P}(\mathcal{S} \times \mathcal{A})$  be probability measures, Lebesgue-absolutely continuous in  $\mathcal{S}$  Define

$$\kappa(m, \nu_1, \nu_2) = \sup_{\pi_1, \dots, \pi_m} \left[ \mathbb{E}_{\nu_2} \left( \frac{d(P^{\pi_m} \dots P^{\pi_1} \nu_1)}{d\nu_2} \right)^2 \right]^{1/2}$$

Let  $\nu$  be the sampling distribution from the algorithm, and  $\mu$  the distribution over which we measure the error in the main theorem, then we assume

$$(1 - \gamma)^2 \sum_{m \geq 1} \gamma^{m-1} m \kappa(m, \mu, \nu) = \phi_{\mu, \nu} < \infty$$